

nVIDIA®

**GPU Computing:
The Democratization of Parallel Computing**

**David Luebke
NVIDIA Research**

Parallel Computing's Golden Age



- **1980s, early `90s: a golden age for parallel computing**
 - Particularly data-parallel computing
- **Architectures**
 - Connection Machine, MasPar, Cray
 - True supercomputers: incredibly exotic, powerful, expensive
- **Algorithms, languages, & programming models**
 - Solved a wide variety of problems
 - Various parallel algorithmic models developed
 - P-RAM, V-RAM, circuit, hypercube, etc.

Parallel Computing's Dark Age



- **But...impact of data-parallel computing limited**
 - Thinking Machines sold 7 CM-1s (100s of systems total)
 - MasPar sold ~200 systems
- **Commercial and research activity subsided**
 - Massively-parallel machines replaced by clusters of ever-more powerful commodity microprocessors
 - Beowulf, Legion, grid computing, ...

Massively parallel computing lost momentum to the inexorable advance of commodity technology

Enter the GPU



- **GPU = *Graphics Processing Unit***
 - Chip in computer video cards, PlayStation 3, Xbox, etc.
 - Two major vendors: NVIDIA and ATI (now AMD)



Enter the GPU



- **GPUs are massively multithreaded manycore chips**
 - **NVIDIA Tesla products have up to 128 scalar processors**
 - **Over 12,000 concurrent threads in flight**
 - **Over 470 GFLOPS sustained performance**
- **Users across science & engineering disciplines are achieving 100x or better speedups on GPUs**
- **CS researchers can use GPUs as a research platform for manycore computing: arch, PL, numeric, ...**

Enter CUDA




- **CUDA** is a scalable parallel programming model and a software environment for parallel computing
 - Minimal extensions to familiar C/C++ environment
 - Heterogeneous serial-parallel programming model
- NVIDIA's **TESLA** GPU architecture accelerates CUDA
 - Expose the computational horsepower of NVIDIA GPUs
 - Enable general-purpose *GPU computing*
- CUDA also maps well to multicore CPUs!

The Democratization of Parallel Computing


- GPU Computing with CUDA brings data-parallel computing to the masses
 - Over 46,000,000 CUDA-capable GPUs sold
 - A “developer kit” costs ~\$200 (for 500 GFLOPS)
- Data-parallel supercomputers are everywhere!
 - CUDA makes this power accessible
 - We’re already seeing innovations in data-parallel computing

Massively parallel computing has become a commodity technology!




nVIDIA®


GPU Computing: Motivation




45X




17X



100X




13–457x



110-240X

GPU Computing: Motivation




35X

GPUs Are Fast



- **Theoretical peak performance: 518 GFLOPS**
- **Sustained μ benchmark performance:**
 - Raw math: 472 GFLOPS (8800 Ultra)
 - Raw bandwidth: 80 GB per second (Tesla C870)
- **Actual application performance:**
 - Molecular dynamics: 290 GFLOPS (VMD ion placement)

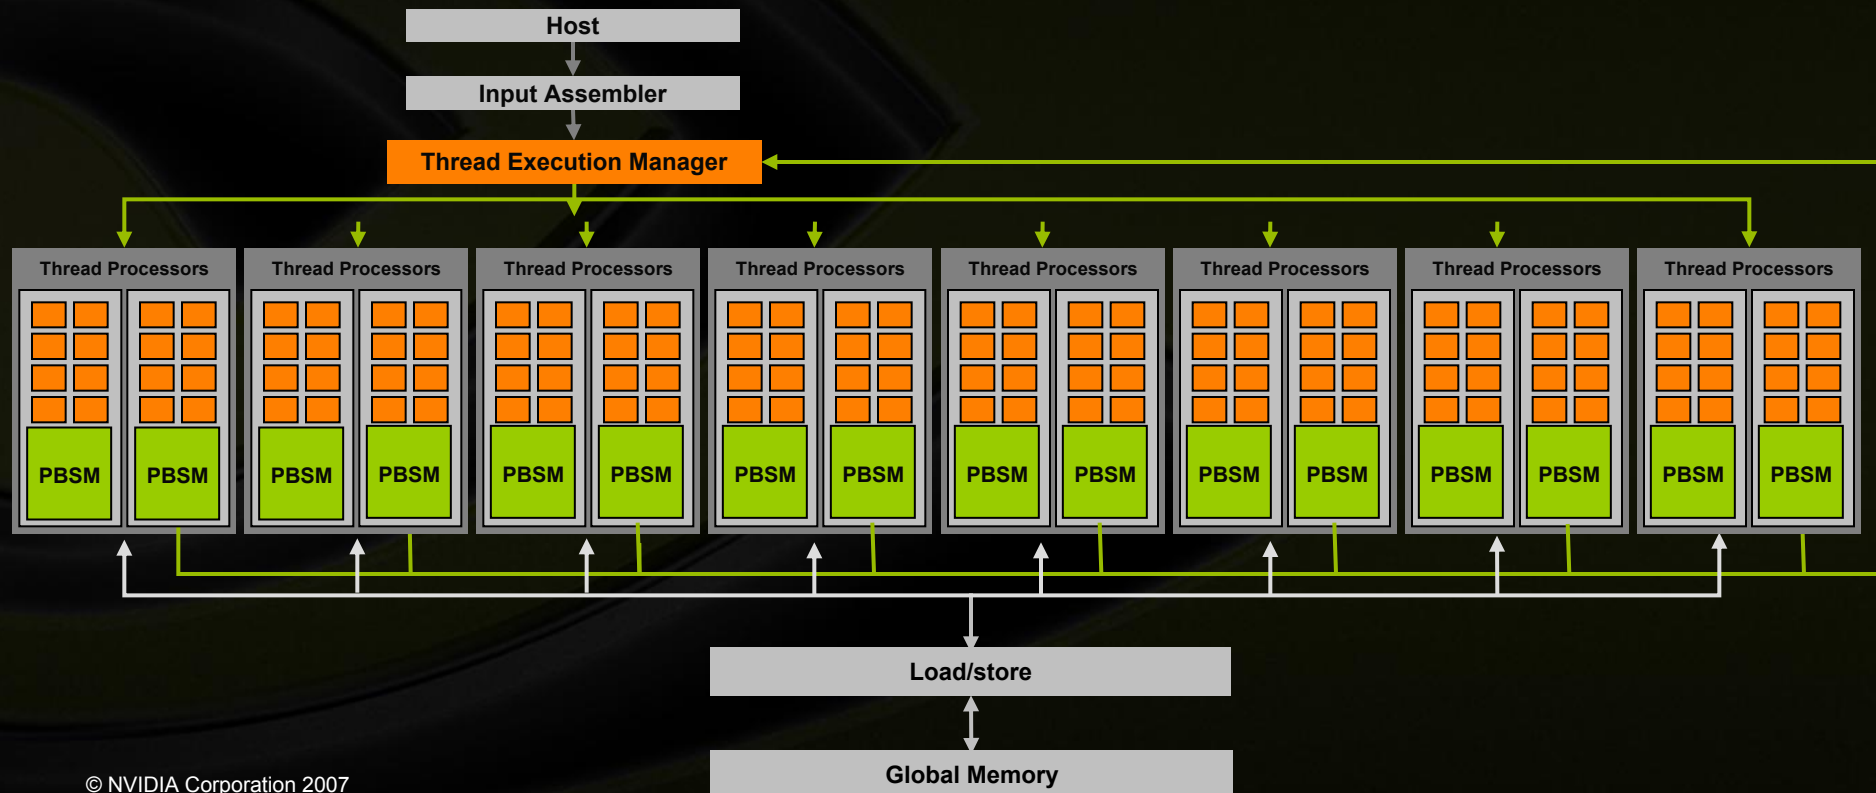
GPUs Are Getting Faster, Faster




Manycore GPU – Block Diagram



- G80 (launched Nov 2006 – GeForce 8800 GTX)
- 128 Thread Processors execute kernel threads
- Up to 12,288 parallel threads active
- Per-block shared memory (PBSM) accelerates processing





nVIDIA®

CUDA Programming Model

Heterogeneous Programming




- **CUDA = serial program with parallel kernels, all in C**
 - Serial C code executes in a CPU thread
 - Parallel kernel C code executes in **thread blocks** across multiple processing elements

Serial Code

Parallel Kernel


```
KernelA<<< nBlk, nTid >>>(args);
```



Serial Code

Parallel Kernel

```
KernelB<<< nBlk, nTid >>>(args);
```



GPU Computing with CUDA: A Highly Multithreaded Coprocessor

- The GPU is a highly parallel **compute device**
 - serves as a coprocessor for the **host** CPU
 - has its own **device memory** on the card
 - executes many **threads** in parallel
- Parallel **kernels** run a single program in many threads
- GPU threads are extremely lightweight
 - Thread creation and context switching are essentially free
- GPU expects 1000's of threads for full utilization

CUDA: Programming GPU in C



- Philosophy: provide minimal set of extensions necessary to expose power

- Declaration specifiers to indicate where things live

```
__global__ void KernelFunc(...); // kernel function, runs on device
__device__ int GlobalVar;         // variable in device memory
__shared__ int SharedVar;         // variable in per-block shared memory
```

- Extend function invocation syntax for parallel kernel launch

```
KernelFunc<<<500, 128>>>(...); // launch 500 blocks w/ 128 threads each
```

- Special variables for thread identification in kernels

```
dim3 threadIdx; dim3 blockIdx; dim3 blockDim; dim3 gridDim;
```

- Intrinsics that expose specific operations in kernel code

```
__syncthreads(); // barrier synchronization within kernel
```

Decoder Ring



GeForce®

Entertainment




Quadro®

Design & Creation



Tesla™

High Performance Computing



GPU

Architecture: TESLA

Chips: G80, G84, G92, ...

A New Platform: Tesla



- **HPC-oriented product line**

- C870: board (1 GPU)
- D870: deskside unit (2 GPUs)
- S870: 1u server unit (4 GPUs)



Conclusion



- **GPUs are massively parallel manycore computers**
 - Ubiquitous - most successful parallel processor in history
 - Useful - users achieve huge speedups on real problems
- **CUDA is a powerful parallel programming model**
 - Heterogeneous - mixed serial-parallel programming
 - Scalable - hierarchical thread execution model
 - Accessible - minimal but expressive changes to C
- **They provide tremendous scope for innovative, impactful research**